

Work

Definition.

The **work** of a program (on a given input) is the sum total of all the operations executed by the program.



pdf element



Precomputation

The idea of **precomputation** is to perform calculations in advance so as to avoid doing them at “mission-critical” times.

Example: Binomial coefficients

$$\binom{n}{k} = \frac{n!}{k!(n-k)!}$$

Computing the “choose” function by implementing this formula can be expensive (lots of multiplications), and watch out for integer overflow for even modest values of **n** and **k**.

Idea: Precompute the table of coefficients when initializing, and perform table look-up at runtime.

Inlining

The idea of **inlining** is to avoid the overhead of a function call by replacing a call to the function with the body of the function itself.

```
double square(double x) {  
    return x*x;  
}  
  
double sum_of_squares(double *A, int n) {  
    double sum = 0.0;  
    for (int i = 0; i < n; ++i) {  
        sum += square(A[i]);  
    }  
    return sum;  
}
```

```
double sum_of_squares(double *A, int n) {  
    double sum = 0.0;  
    for (int i = 0; i < n; ++i) {  
        double temp = A[i];  
        sum += temp*temp;  
    }  
    return sum;  
}
```

Packing and Encoding

The idea of **packing** is to store more than one data value in a machine word. The related idea of **encoding** is to convert data values into a representation requiring fewer bits.

Example: Encoding dates

- The string “September 11, 2018” can be stored in 18 bytes — more than two double (64-bit) words — which must be moved whenever a date is manipulated.
- Assuming that we only store years between 4096 B.C.E. and 4096 C.E., there are about $365.25 \times 8192 \approx 3\text{ M}$ dates, which can be encoded in $\lceil \lg(3 \times 10^6) \rceil = 22$ bits, easily fitting in a single (32-bit) word.
- But determining the month of a date takes more work than with the string representation.

Why Learn Bit Hacks?

Why learn bit hacks if they don't even work?

- Because the compiler does them, and it will help to understand what the compiler is doing when you look at the assembly code.
- Because sometimes the compiler doesn't optimize, and you have to do it yourself by hand.
- Because many bit hacks for words extend naturally to bit and word hacks for vectors.
- Because these tricks arise in other domains, and so it pays to be educated about them.
- Because they're fun!

Set the kth Bit

Problem

Set k th bit in a word x to 1.

Idea

Shift and OR.

$y = x | (1 << k);$

Example

$k = 7$

x	1011110101101101
$1 << k$	0000000010000000
$x (1 << k)$	1011110111101101

Clear the kth Bit

Problem

Clear the k th bit in a word x .

Idea

Shift, complement, and AND.

```
y = x & ~(1 << k);
```

Example

$k = 7$

x	101111011101101
$1 << k$	000000010000000
$\sim(1 << k)$	111111101111111
$x \& \sim(1 << k)$	1011110101101101

Toggle the kth Bit

Problem

Flip the k th bit in a word x .

Idea

Shift and XOR.

$y = x \wedge (1 \ll k);$

Example ($0 \rightarrow 1$)

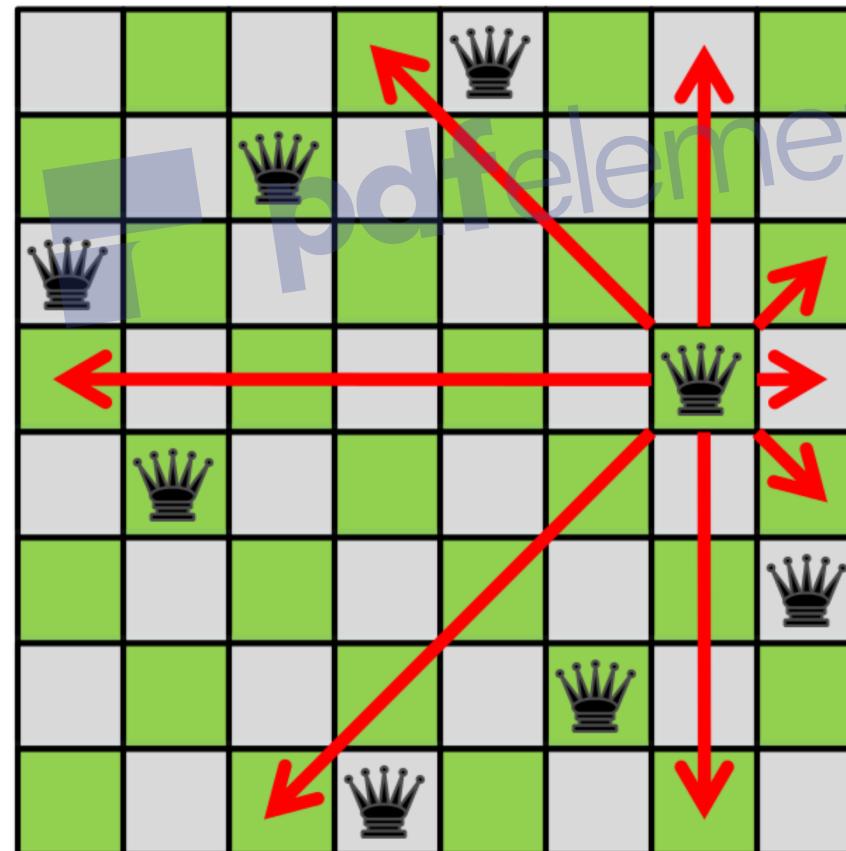
$k = 7$

x	1011110101101101
$1 \ll k$	0000000010000000
$x \wedge (1 \ll k)$	1011110111101101

Queens Problem

Problem

Place n queens on an $n \times n$ chessboard so that no queen attacks another, i.e., no two queens in any row, column, or diagonal. Count the number of possible solutions.

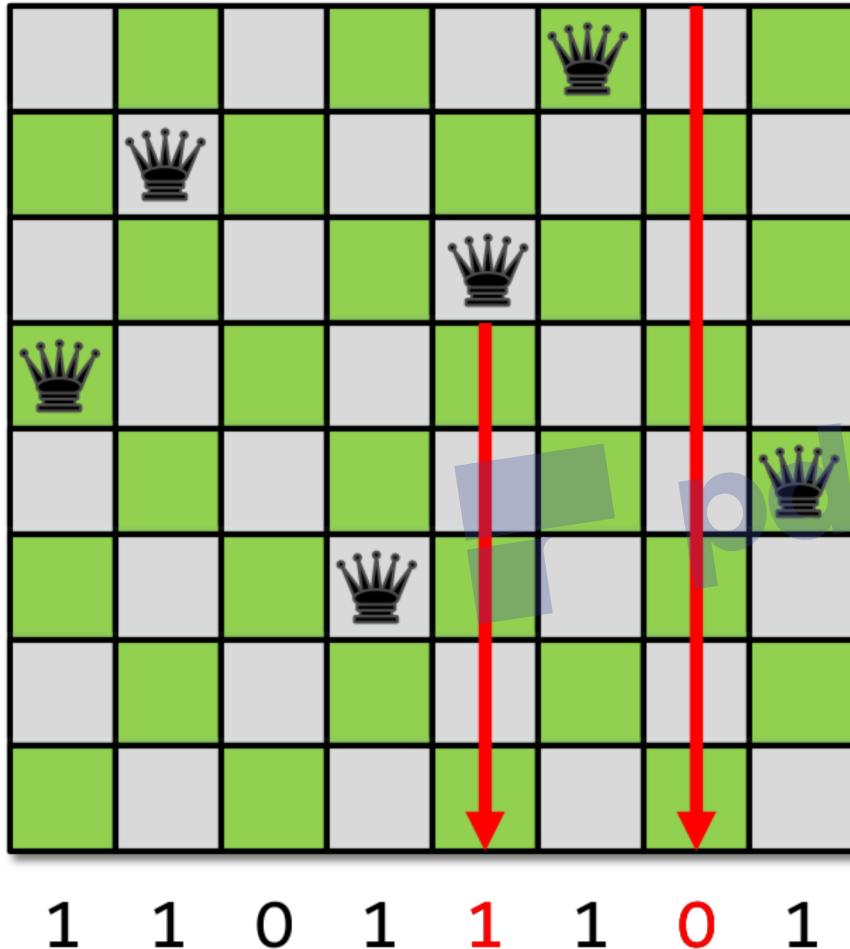


Board Representation

The backtrack search can be implemented as a simple recursive procedure, but how should the board be represented to facilitate queen placement?

- array of n^2 bytes?
- array of n^2 bits?
- array of n bytes?
- 3 bitvectors of size n , $2n-1$, and $2n-1$.

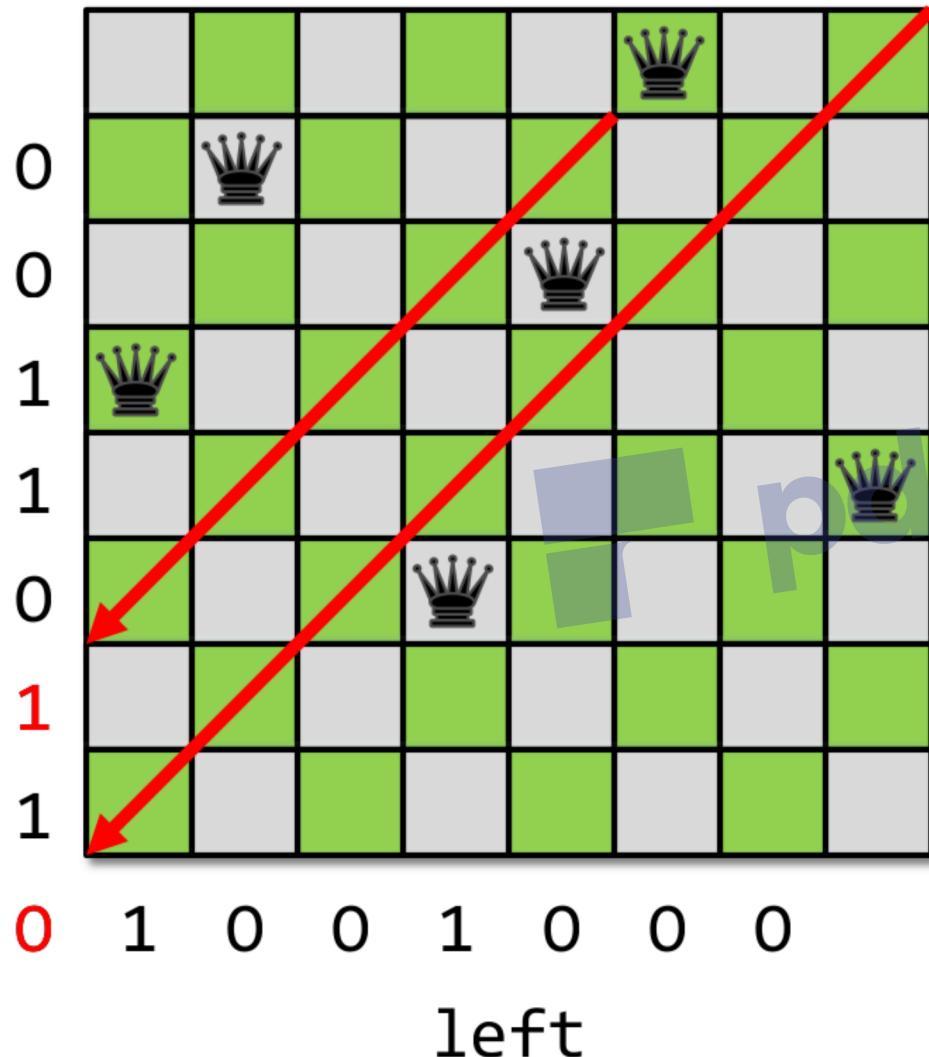
Bitvector Representation



Placing a queen in column c is not safe if
down & ($1 \ll c$);
is nonzero.

element

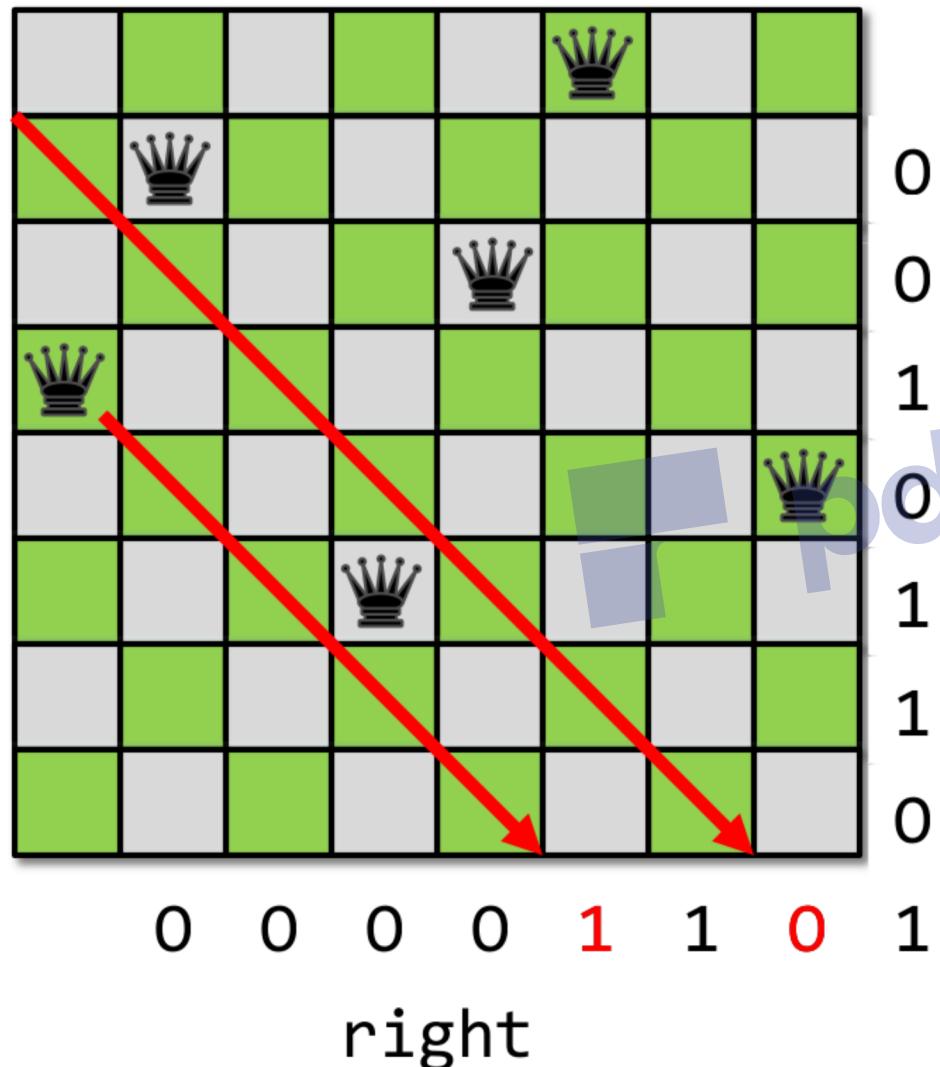
Bitvector Representation



Placing a queen in row r and column c is not safe if

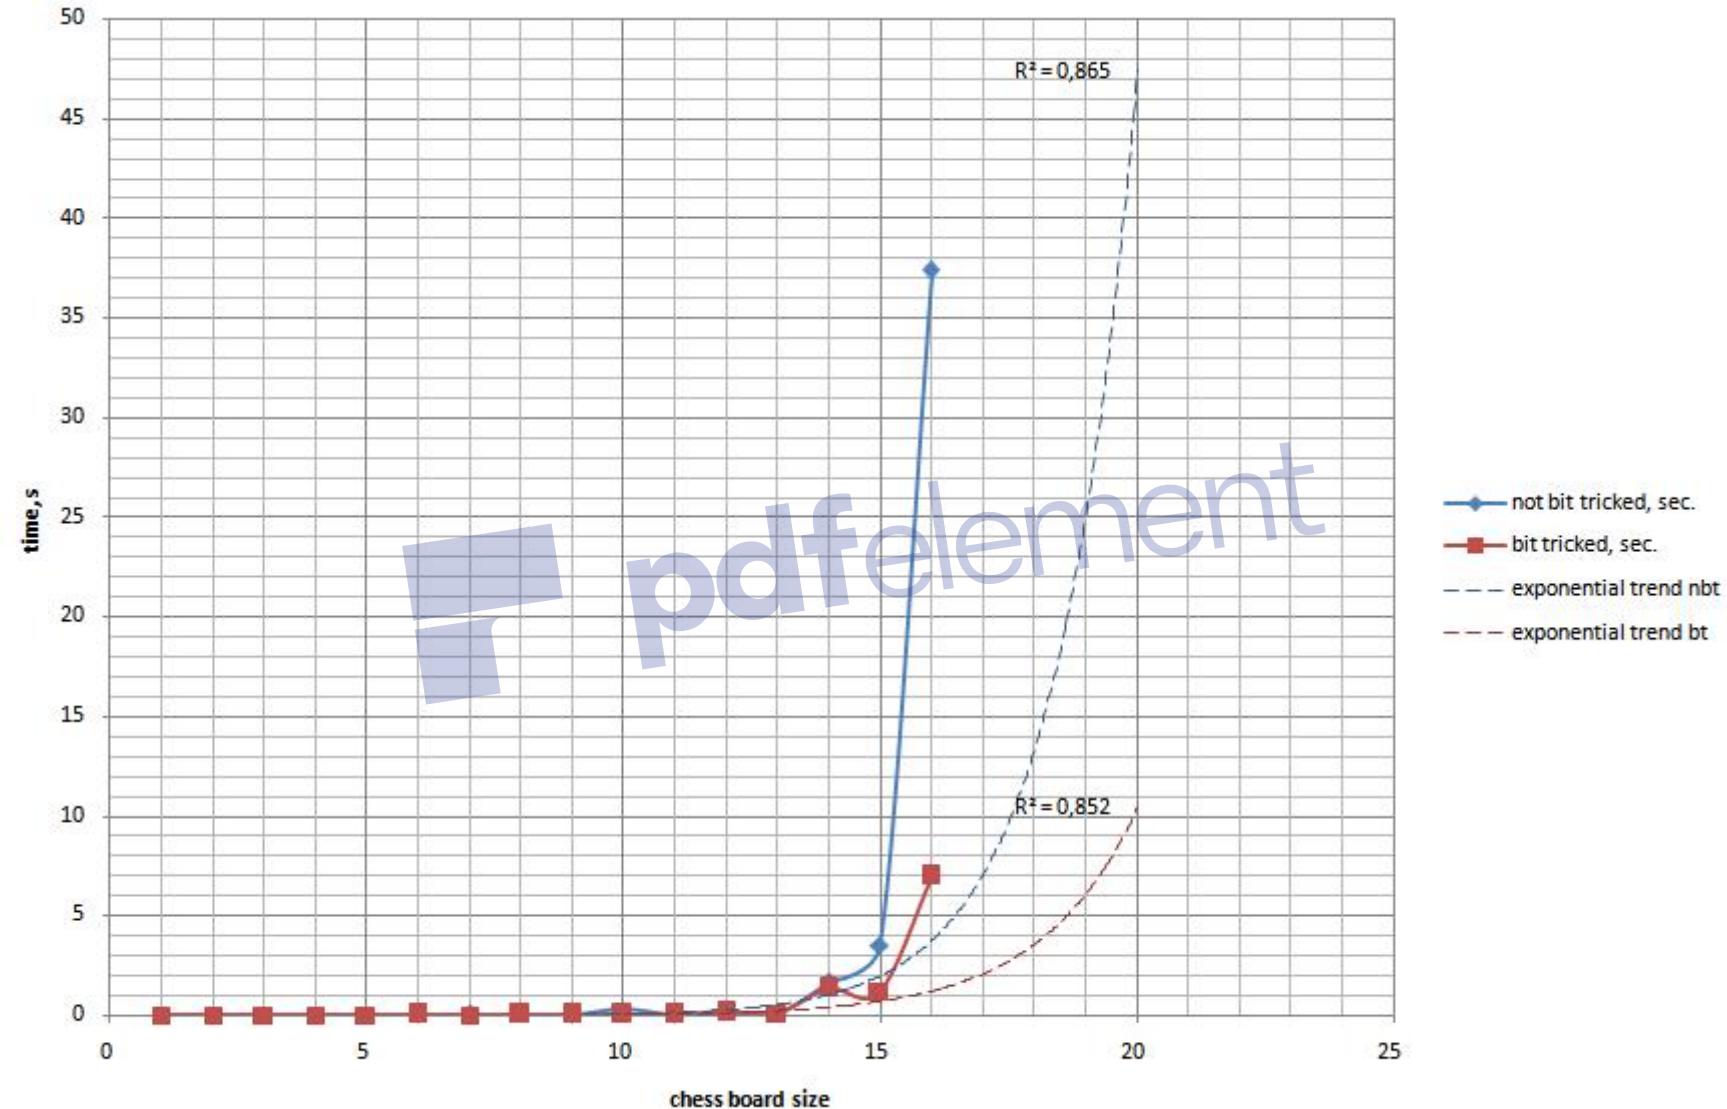
$\text{left} \& (1 << (r+c))$ is nonzero.

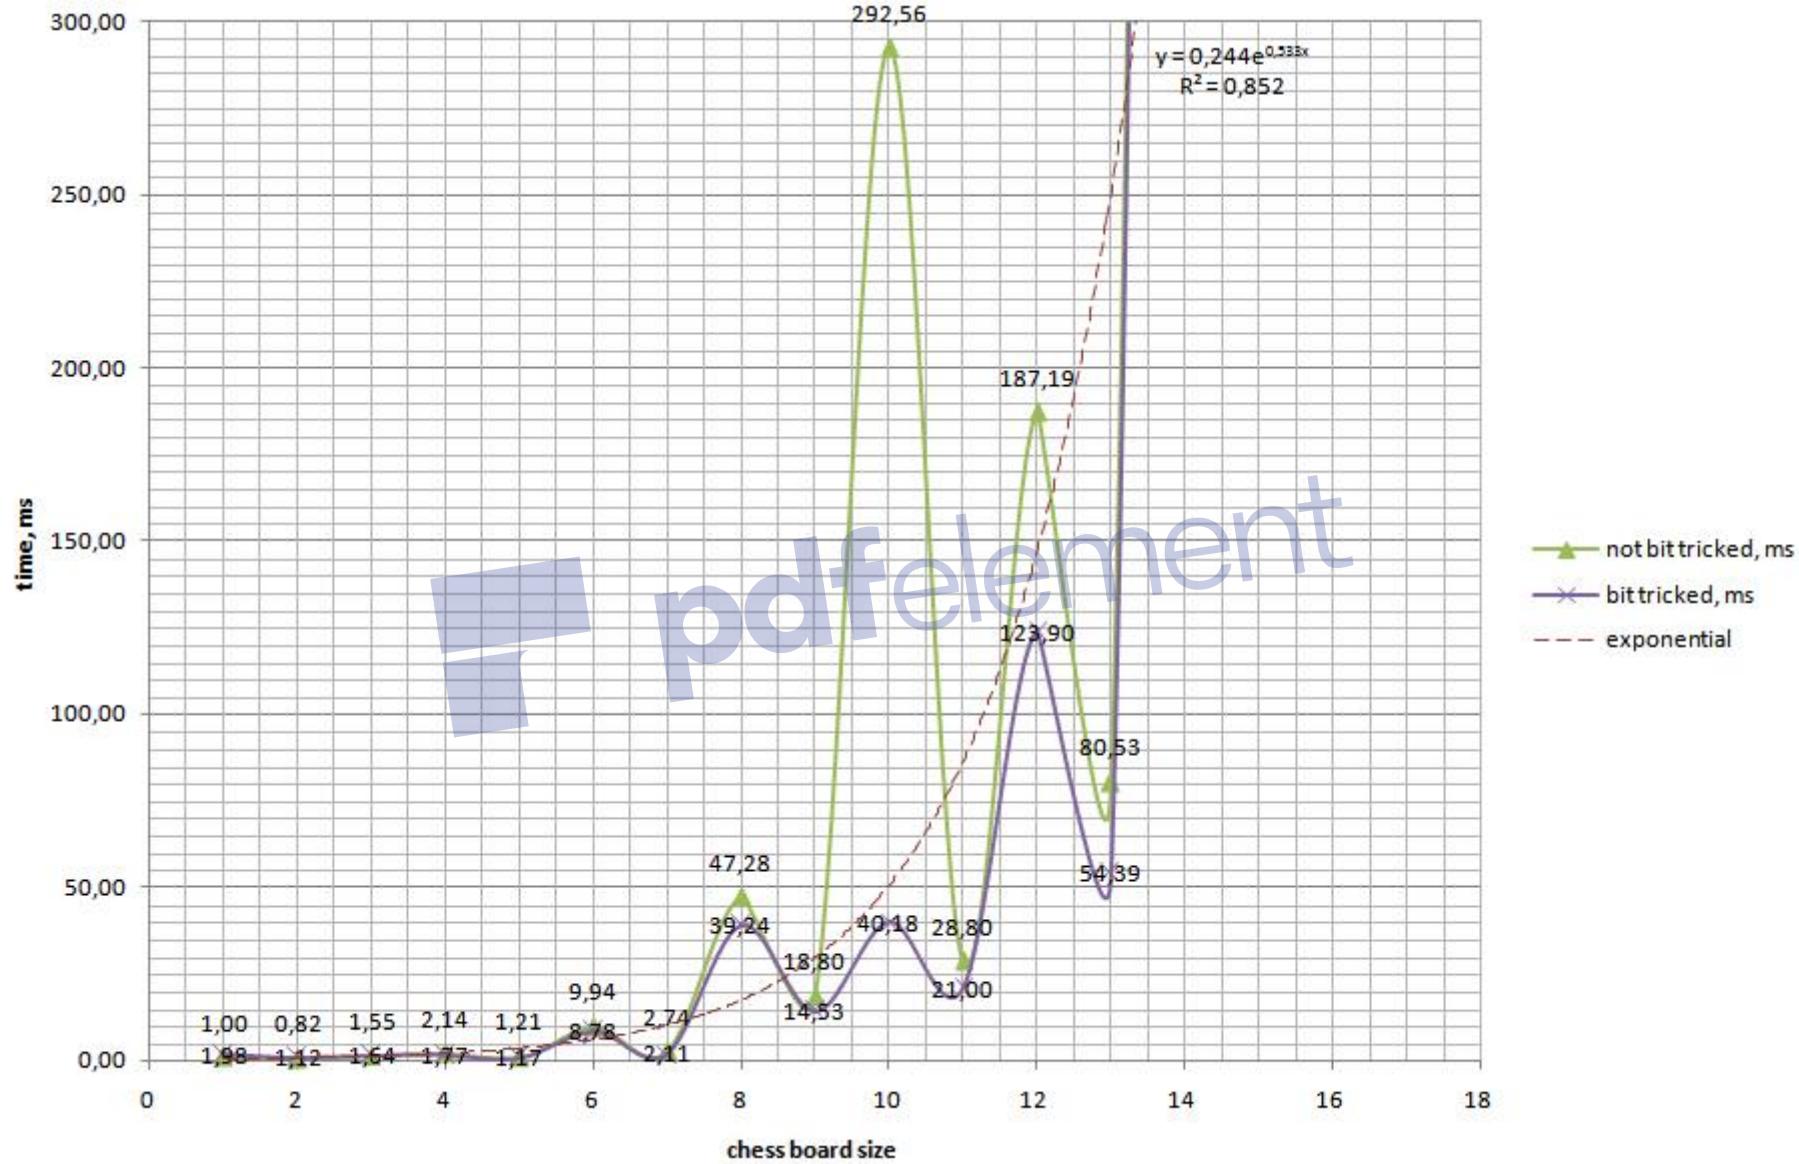
Bitvector Representation



Placing a queen in row r and column c is not safe if
 $\text{right} \& (1 << (n-1-r+c))$
is nonzero.

pdf element





Dynamic Frequency and Voltage Scaling

DVFS is a technique to reduce power by **adjusting** the **clock frequency** and **supply voltage** to transistors.

- Reduce operating frequency if chip is too hot or otherwise to conserve (especially battery) power.
- Reduce voltage if frequency is reduced.

$$\text{Power} \propto C V^2 f$$

C = dynamic capacitance

≈ roughly area × activity (how many bits toggle)

V = supply voltage

f = clock frequency

Reducing frequency and voltage results in a cubic reduction in power (and heat).

But it wreaks havoc on performance measurements!

Sources of Variability

- Daemons and background jobs
- Interrupts
- Code and data alignment
- Thread placement
- Runtime scheduler
- Hyperthreading
- Multitenancy
- Dynamic voltage and frequency scaling (DVFS)
- Turbo Boost
- Network traffic